

Question 1

Please see `deployment.yaml` and `service.yaml`

Following is the output of some of the checker commands for your convenience

- `kubectl get all`

```
NAME                                         READY   STATUS    RESTARTS   AGE
pod/rabbitmq-deployment-7ff8c859c6-2qwhr   1/1     Running   0           69m
pod/rabbitmq-deployment-7ff8c859c6-p94bl    1/1     Running   0           69m
pod/rabbitmq-deployment-7ff8c859c6-vh7d2    1/1     Running   0           69m
pod/rabbitmq-deployment-7ff8c859c6-zgqxg    1/1     Running   0           69m
pod/rabbitmq-deployment-7ff8c859c6-zxh8r    1/1     Running   0           69m

NAME                TYPE          CLUSTER-IP      EXTERNAL-IP      PORT(S)
AGE
service/kubernetes  ClusterIP     10.152.183.1    <none>           443/TCP
32h
service/rabbitmq-svc LoadBalancer  10.152.183.208  <pending>
32500:32500/TCP    70m

NAME                                         READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/rabbitmq-deployment         5/5      5             5           69m

NAME                                         DESIRED   CURRENT   READY
AGE
replicaset.apps/rabbitmq-deployment-7ff8c859c6 5         5         5
69m
```

- `kubectl describe deployment.apps/rabbitmq-deployment`

```
Name:                rabbitmq-deployment
Namespace:            default
CreationTimestamp:    Wed, 16 Oct 2019 18:25:56 -0500
Labels:               <none>
Annotations:          deployment.kubernetes.io/revision: 1
                     kubect1.kubernetes.io/last-applied-configuration:

{"apiVersion":"apps/v1","kind":"Deployment","metadata":{"annotations":
{},"name":"rabbitmq-deployment","namespace":"default"},"spec":{"repl...
Selector:             app=rabbit
Replicas:              5 desired | 5 updated | 5 total | 5 available | 0
unavailable
```

```

StrategyType:      RollingUpdate
MinReadySeconds:    0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=rabbit
  Containers:
    rabbitmq:
      Image:      rabbitmq:3-management
      Port:        15672/TCP
      Host Port:   0/TCP
      Limits:
        cpu:      125m
        memory:    300Mi
      Requests:
        cpu:        50m
        memory:      150Mi
      Environment:  <none>
      Mounts:        <none>
      Volumes:        <none>
  Conditions:
    Type           Status  Reason
    ----           -
    Available       True    MinimumReplicasAvailable
    Progressing     True    NewReplicaSetAvailable
OldReplicaSets:  <none>
NewReplicaSet:   rabbitmq-deployment-7ff8c859c6 (5/5 replicas created)
Events:          <none>

```

- `kubectl describe service/rabbitmq-svc`

```

Name:      rabbitmq-svc
Namespace: default
Labels:    app=rabbit
Annotations: kubectrl.kubernetes.io/last-applied-configuration:
              {"apiVersion":"v1","kind":"Service","metadata":
{"annotations":{},"labels":{"app":"rabbit"},"name":"rabbitmq-
svc","namespace":"default"},"s...
Selector:  app=nginx
Type:      LoadBalancer
IP:        10.152.183.208
Port:      <unset> 32500/TCP
TargetPort: 15672/TCP
NodePort:   <unset> 32500/TCP
Endpoints:  <none>

```

Session Affinity:	None
External Traffic Policy:	Cluster
Events:	<none>

Question 2

For the given `interesting.pcap` file, I noticed that there are a huge amount of [SYN] signals sent to the same target from different IP addresses. My assumption is someone is conducting a DOS attack to the server, which aims to block the three way handshake from other clients.

Question 3

```
policy_module(bank, 1.0.0)
type bank_exec_t;
type bank_admin_t;
type bank_balance_t;
init_daemon_domain(bank_admin_t, bank_exec_t)
files_type(bank_balance_t)
require {
    type bank_admin_t;
    type bank_balance_t;
    type user_home_t;
    class file { create execute open read write getattr };
    class dir { add_name search write create remove_name getattr };
}
type_transition initrc_t bank_exec_t:file bank_exec_t;
allow user_home_t bank_balance_t:dir { read getattr };
allow bank_admin_t bank_balance_t:file { read write delete create getattr };
allow bank_admin_t bank_balance_t:dir { create getattr };
allow user_home_t user_home_t:dir { create open read write getattr };
```

Question 4

- First I used this command to find the syscalls that is related to `/var/www`

```
sudo sysdig -r suspicious.scap | grep /var/www
```

- Then I located the PIDs of the process that makes calls to `/var/www`, which is `18886, 18888, 18864, 18865, 18866`. Then I used the following command to see the exact syscalls that were made by that process:
- In order to make sure the above processes are actually making large network bandwidth usage, I ran the following command: `sysdig -r suspicious.scap -c topprocs_net`, which gave me the following result: (* are the pids that we found suspicious)

Bytes	Process	PID
-------	---------	-----

```
-----
---
634.46KB      kube-apiserver      6314
240.79KB      etcd                 6007
176.45KB      heapster             9798
139.40KB      influxd              9965
131.46KB      kube-controller      6605
64.88KB       kubelet              6349
46.69KB       coredns              8870
37.94KB       kube-scheduler       6603
35.25KB       pod_nanny            10647
35.25KB       pod_nanny            10442
20.54KB       kube-proxy           6320
18.77KB       Socket               7179
13.17KB       systemd-resolve      452
12.17KB       dashboard            10031
9.07KB        apache2              18866 *
3.85KB        apache2              18864 *
3.85KB        apache2              18888 *
3.14KB        apache2              18886 *
2.98KB        DNS                  7179
1.47KB        mysqld               19358
931B          apache2              18859
804B          apache2              18865 *
395B          eventer              10193
235B          NetworkManager       636
```

- Finally, in order to find out where the traffic is coming from, I ran the following command to check the IP addresses of each process (I used system call `accept()` as filter because it specifies the source IP and target IP as its inputs:

```
sudo sysdig -r suspicious.scap \
proc.pid=18886 or \
proc.pid=18888 or \
proc.pid=18864 or \
proc.pid=18865 or \
proc.pid=18866 \
| grep accept
```

The above script generates the following result:

```

174425 17:51:08.958745443 0 apache2 (18886) < accept
fd=10(<4t>10.1.1.1:43692->10.1.1.28:80) tuple=10.1.1.1:43692->10.1.1.28:80
queuepct=0 queuelen=0 queuemax=128
434940 17:51:13.971452480 0 apache2 (18886) > accept flags=0
460020 17:51:14.726287385 1 apache2 (18888) < accept
fd=10(<4t>10.1.1.1:43712->10.1.1.28:80) tuple=10.1.1.1:43712->10.1.1.28:80
queuepct=0 queuelen=0 queuemax=128
663023 17:51:19.734840292 1 apache2 (18888) > accept flags=0
987491 17:51:28.024501242 0 apache2 (18864) < accept
fd=10(<4t>10.1.1.1:43762->10.1.1.29:80) tuple=10.1.1.1:43762->10.1.1.29:80
queuepct=0 queuelen=0 queuemax=128
1198435 17:51:33.032192827 0 apache2 (18864) > accept flags=0
1229246 17:51:33.537011565 0 apache2 (18865) < accept
fd=10(<4t>10.1.1.1:43780->10.1.1.29:80) tuple=10.1.1.1:43780->10.1.1.29:80
queuepct=0 queuelen=0 queuemax=128
1424070 17:51:38.543233202 0 apache2 (18865) > accept flags=0
1522863 17:51:40.930099309 0 apache2 (18866) < accept
fd=10(<4t>10.1.1.1:43804->10.1.1.29:80) tuple=10.1.1.1:43804->10.1.1.29:80
queuepct=0 queuelen=0 queuemax=128

```

And we can see that the traffic all comes from this same IP address: `10.1.1.1`

In order to get more information, I ran `sysdig-inspect` and used the following filter command:

```

proc.pid=18886 or proc.pid=18888 or proc.pid=18864 or proc.pid=18865 or
proc.pid=18866

```

PID	VPID	CPU USER	TH	VIRT	RES	FILE	NET	CONTAINER	Command
18888	29	0 www-data	1	85848064	15392768	2253	3949		/usr/sbin/apache2 -k start
18866	30	0 www-data	1	85975040	15884288	12596	9290		/usr/sbin/apache2 -k start
18886	27	0 www-data	1	85966848	15511552	2159	3223		/usr/sbin/apache2 -k start
18864	28	0 www-data	1	85966848	15810560	2253	3949		/usr/sbin/apache2 -k start
18865	29	0 www-data	1	85975040	15683584	675	806		/usr/sbin/apache2 -k start